

PageRank

Une initialisation avec un PR de $\text{damping_factor} / \text{total_noeuds}$ sur tous les noeuds

Une itérations du calcul du PR, il suffit d'itérer sur le job CalculPR

| noeud\itération | 1 | 10 | 20 | 50 | 75 | 100 |
|-----------------|--------|-------|-------|-------|-------|-------|
| 0 | 0,1503 | 38,19 | 48,54 | 50,65 | 50,66 | 50,66 |
| 1 | 0,1505 | 48,45 | 60,77 | 63,28 | 63,29 | 63,29 |
| 10 | 0,1501 | 17,88 | 23,74 | 24,95 | 24,96 | 24,96 |
| 100 | 0,1501 | 18,94 | 25,23 | 26,53 | 26,53 | 26,53 |
| 1000 | 0,1500 | 1,584 | 2,041 | 2,136 | 2,136 | 2,136 |

On constate qu'au bout de 20 itérations les scores du PageRank s'approchent fortement du score final. Après 50 itérations le score ne varie presque plus : ex entre 75 et 100 pas de variations

```
1 package cs.bigdata.ex52;
2
3 import org.apache.hadoop.conf.Configuration;
4
5 public class PageRank {
6     static int TOTAL_NOEUDS = 75879;
7     static Double DAMPLING_FACTOR = 0.85;
8     static int nIter = 100;
9
10    public static void main(String[] args) throws Exception {
11        // Connection a HDFS
12        Configuration conf = new Configuration();
13        conf.set("fs.defaultFS", "hdfs://localhost:9000");
14
15        // Step 1- Init info
16        InitPR.main(conf);
17
18        // Step 2 - Calcul du PageRank
19
20        for (int i = 1; i <= nIter; i++)
21        {
22            CalculPR.main(conf, i);
23        }
24    }
25 }
```

Console [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_152.jdk/Contents/Home/bin/java (15 janv. 2018 à 13:04)

```
terminated> PageRank [Java Application] /Library/Java/JavaVirtualMachines/jdk1.8.0_152.jdk/Contents/Home/bin/java (15 janv. 2018 à 13:04)
HDFS: Number of large read operations=0
HDFS: Number of write operations=1204
Map-Reduce Framework
Map input records=75880
Map output records=600256
Map output bytes=17078177
Map output materialized bytes=18285308
Input split bytes=137
Combine input records=0
Combine output records=0
Reduce input groups=75880
Reduce shuffle bytes=0
Reduce input records=600256
Reduce output records=75880
Spilled Records=1200512
Shuffled Maps =0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=4
Total committed heap usage (bytes)=3827302400
File Input Format Counters
Bytes Read=4473596
File Output Format Counters
Bytes Written=4472721
```

OMA/PLP/projects master X

► hdfs dfs -cat pageRankOutput21/part-r-00000 | head -n 6

WARNING: An illegal reflective access operation has occurred

WARNING: Illegal reflective access by org.apache.hadoop.security.krb5.Config.getInstance()

WARNING: Please consider reporting this to the maintainers of c

WARNING: Use --illegal-access=warn to enable warnings of further

WARNING: All illegal access operations will be denied in a future

18/01/15 13:16:03 WARN util.NativeCodeLoader: Unable to load native

19925.5048393553

0 48.53697199108769 39,40,41,42,43,44,45,46,47,48,4

6,97,98,99,100,101,102,103,38,37,36,4,104,105,106,107,108,109,1

124,125,126,127,128,129,130,131,132,133,134,135,136,137,138,139

1 60.77209580068582 24,21,20,17,16,11,10,5,4,3,1988

02,2470,2464,2256,2144,2065,2046,2040,2036,1778,1747,1718,1642,

5,728,727,725,723,704,694,689,684,663,649,645,625,592,586,576,5

7,426,425,424,423,422,421,420,419,418,417,416,415,414,413,412,4

5,374,373,372,371,370,369,368,367,366,365,364,363,362,361,360,3

3,322,321,320,319,318,317,316,315,314,313,312,311,310,300,297,2

2,91,90,87,86,85,83,82,81,78,77,75,74,71,68,67,66,65,63,61,58,5

10 23.740468998782667 48,50,51,61,65,66,69,72,73,74,7

426,423,426,440,545,540,550,555,550,586,611,612,620,625,627,667